

HPC - Lab 03

Aubry Mangold
ISCS, HEIG-VD
aubry.mangold@heig-vd.ch

Compilation optimizations

Introduction

The following report presents a series of optimizations applied to C code, comparing naive implementations with manually optimized and compiler-optimized versions. The goal is to demonstrate how different optimization techniques can improve performance and reduce code complexity.

Case 1 - Recursion to iteration

A recursive function can be transformed into an iterative one by using a loop and variables. This modification removes the overhead of function calls and stack management.

<pre>1 #pragma GCC push_options 2 #pragma GCC optimize("-O0") 3 int factorial(int n) { 4 if (n <= 1) return 1; 5 return n * factorial(n - 1); 6 } 7 #pragma GCC pop_options</pre>	<pre>1 #pragma GCC push_options 2 #pragma GCC optimize("-O0") 3 int factorial_manual(int n) { 4 int result = 1; 5 for (int i = 2; i <= n; ++i) { 6 result *= i; 7 } 8 return result; 9 } 10 #pragma GCC pop_options</pre>	<pre>1 #pragma GCC push_options 2 #pragma GCC optimize("O2") 3 int factorial_compiler(int n) { 4 if (n <= 1) return 1; 5 return n * factorial_compiler(n 6 - 1); 7 } 8 #pragma GCC pop_options</pre>
--	--	---

Table 1: The naive, manually-optimized, and compiler-optimized versions of the factorial function.

The naive version is a simple factorial function that uses recursion. The manually optimized version replaces the recursive calls with a loop, which is more efficient in terms of stack usage and function call overhead. The compiler-optimized version uses the same recursive logic but applies compiler optimizations to improve performance. It wasn't clear which flags applied the optimisation in a discrete manner, so O2 was used.

<pre>1 factorial: 2 push rbp 3 mov rbp, rsp 4 sub rsp, 16 5 mov DWORD PTR [rbp-4], edi 6 cmp DWORD PTR [rbp-4], 1 7 jg .L2 8 mov eax, 1 9 jmp .L3 10 .L2: 11 mov eax, DWORD PTR [rbp-4] 12 sub eax, 1 13 mov edi, eax 14 call factorial 15 imul eax, DWORD PTR [rbp-4] 16 .L3: 17 leave 18 ret</pre>	<pre>1 factorial_manual: 2 push rbp 3 mov rbp, rsp 4 mov DWORD PTR [rbp-20], edi 5 mov DWORD PTR [rbp-4], 1 6 mov DWORD PTR [rbp-8], 2 7 jmp .L5 8 .L6: 9 mov eax, DWORD PTR [rbp-4] 10 imul eax, DWORD PTR [rbp-8] 11 mov DWORD PTR [rbp-4], eax 12 add DWORD PTR [rbp-8], 1 13 .L5: 14 mov eax, DWORD PTR [rbp-8] 15 cmp eax, DWORD PTR [rbp-20] 16 jle .L6 17 mov eax, DWORD PTR [rbp-4] 18 pop rbp 19 ret</pre>	<pre>1 factorial_compiler: 2 mov eax, 1 3 cmp edi, 1 4 jle .L8 5 .L9: 6 mov edx, edi 7 sub edi, 1 8 imul eax, edx 9 cmp edi, 1 10 jne .L9 11 .L8: 12 ret</pre>
---	---	---

Table 2: Assembly traductions of the factorial functions using gcc 14.2 x86-64.

Inspecting the assembly displayed in Table 2, we observe that the unoptimized version uses the stack and contains an actual function call to itself on line 14. It uses a lot of time and space on overhead.

The manually optimized version of the function uses manual iteration while saving values on the stack, although it still does a lot of moving and comparisons. In particular, it has no nested calls, which removes the penalty for maintaining a multi-level stack frame.

The compiler optimized version using O2 as a flag is much more brief and efficient, as it transformed the factorial function into a loop but was also able to reduce the number of instruction (particularly the number of mov instructions) quite a bit.

Case 2 - Branch removal

<pre>1 #pragma GCC push_options 2 #pragma GCC optimize("-O0") 3 int branch(int x) { 4 if (x == 0) 5 return 1; 6 return 0; 7 } 8 #pragma GCC pop_options</pre>	<pre>1 #pragma GCC push_options 2 #pragma GCC optimize("-O0") 3 int branch_manual(int x) { 4 return x == 0; 5 } 6 #pragma GCC pop_options</pre>	<pre>1 #pragma GCC push_options 2 #pragma GCC optimize("O2") 3 int branch_compiler(int x) { 4 if (x == 0) 5 return 1; 6 return 0; 7 } 8 #pragma GCC pop_options</pre>
---	---	---

Table 3: The naive, manually-optimized, and compiler-optimized versions of the branch removal function.

Branch removal is a technique used to eliminate unnecessary branches in code, which can in turn improve performance by reducing the number of conditional checks and branch prediction misses. The code displayed in Table 3 shows the naive, manually optimized, and compiler-optimized versions of a function that checks if an integer is equal to zero. The naive version, the function uses an if statement to check if the input is zero and returns 1 if true, otherwise it returns 0. The manually optimized version simplifies this logic by directly returning the result of the comparison ($x == 0$). The compiler-optimized version uses the same logic as the naive version but applies compiler optimizations O2 to apply branch removal. If seemed the if-conversion and if-conversion2 optimizations were not enough to have the compiler remove the check.

<pre>1 branch: 2 push rbp 3 mov rbp, rsp 4 mov DWORD PTR [rbp-4], edi 5 cmp DWORD PTR [rbp-4], 0 6 jne .L2 7 mov eax, 1 8 jmp .L3 9 .L2: 10 mov eax, 0 11 .L3: 12 pop rbp 13 ret</pre>	<pre>1 branch_manual: 2 push rbp 3 mov rbp, rsp 4 mov DWORD PTR [rbp-4], edi 5 cmp DWORD PTR [rbp-4], 0 6 sete al 7 movzx eax, al 8 pop rbp 9 ret</pre>	<pre>1 branch_compiler: 2 xor eax, eax 3 test edi, edi 4 sete al 5 ret</pre>
---	---	--

Table 4: Assembly translations of the branch functions using gcc 14.2 x86-64.

As shown in the assembly displayed in Table 4, the initial version uses a conditional jump call, which may cause branch mispredictions.

The manually optimized version uses branchless logic by directly returning the result of the operation using sete and movzx to zero-extend the result before using it as a return value. This avoids the conditional jump and improves performance by avoiding branch mispredictions, hence increasing pipeline efficiency.

The compiler-optimized version of the function uses the same logic as the manually optimized version but applies compiler optimizations to further improve performance. It uses xor to zero out the register before using test and sete to set the return value based on the comparison. This is a common GCC trick for return value handling, and results in a very concise branchless implementation.

Case 3 - Unswitching

<pre> 1 #pragma GCC push_options 2 #pragma GCC optimize("-O0") 3 void unswitch(int *arr, int flag) 4 { 5 for (int i = 0; i < 3; i++) { 6 if (flag) { 7 arr[i] *= 2; 8 } 9 } 10 #pragma GCC pop_options </pre>	<pre> 1 #pragma GCC push_options 2 #pragma GCC optimize("-O0") 3 void unswitch_manual(int *arr, 4 int flag) { 5 if (flag) { 6 for (int i = 0; i < 3; i 7 +) { 8 arr[i] *= 2; 9 } 10 } 11 #pragma GCC pop_options </pre>	<pre> 1 #pragma GCC push_options 2 #pragma GCC optimize("-O1", "- 3 funswitch-loops") 4 void unswitch_compiler(int *arr, 5 int flag) { 6 for (int i = 0; i < 3; i++) { 7 if (flag) { 8 arr[i] *= 2; 9 } 10 } 11 #pragma GCC pop_options </pre>
---	---	--

Table 5: The naive, manually-optimized, and compiler-optimized versions of the branch removal function.

Unswitching is an optimization whereby a condition is moved outside of a loop to reduce the number of conditional checks performed during each iteration. Such a condition is said to be *loop invariant*. This can improve performance by reducing the number of branches taken in the loop. The naive version does many checks inside the loop, while the manually optimized version moves the check outside of the loop. The compiler-optimized version uses both the `O1` and `-funswitch-loops` options together to get the desired output.

<pre> 1 unswitch: 2 push rbp 3 mov rbp, rsp 4 mov QWORD PTR [rbp-24], rdi 5 mov DWORD PTR [rbp-28], esi 6 mov DWORD PTR [rbp-4], 0 7 jmp .L2 8 .L4: 9 cmp DWORD PTR [rbp-28], 0 10 je .L3 11 mov eax, DWORD PTR [rbp-4] 12 cdqe 13 lea rdx, [0+rax*4] 14 mov rax, QWORD PTR [rbp-24] 15 add rax, rdx 16 mov edx, DWORD PTR [rax] 17 mov eax, DWORD PTR [rbp-4] 18 cdqe 19 lea rcx, [0+rax*4] 20 mov rax, QWORD PTR [rbp-24] 21 add rax, rcx 22 add edx, edx 23 mov DWORD PTR [rax], edx 24 .L3: 25 add DWORD PTR [rbp-4], 1 26 .L2: 27 cmp DWORD PTR [rbp-4], 2 28 jle .L4 29 nop 30 nop 31 pop rbp 32 ret </pre>	<pre> 1 unswitch_manual: 2 push rbp 3 mov rbp, rsp 4 mov QWORD PTR [rbp-24], rdi 5 mov DWORD PTR [rbp-28], esi 6 cmp DWORD PTR [rbp-28], 0 7 je .L9 8 mov DWORD PTR [rbp-4], 0 9 jmp .L7 10 .L8: 11 mov eax, DWORD PTR [rbp-4] 12 cdqe 13 lea rdx, [0+rax*4] 14 mov rax, QWORD PTR [rbp-24] 15 add rax, rdx 16 mov edx, DWORD PTR [rax] 17 mov eax, DWORD PTR [rbp-4] 18 cdqe 19 lea rcx, [0+rax*4] 20 mov rax, QWORD PTR [rbp-24] 21 add rax, rcx 22 add edx, edx 23 mov DWORD PTR [rax], edx 24 add DWORD PTR [rbp-4], 1 25 .L7: 26 cmp DWORD PTR [rbp-4], 2 27 jle .L8 28 .L9: 29 nop 30 pop rbp 31 ret </pre>	<pre> 1 unswitch_compiler: 2 test esi, esi 3 je .L10 4 sal DWORD PTR [rdi] 5 sal DWORD PTR [rdi+4] 6 sal DWORD PTR [rdi+8] 7 .L10: 8 ret </pre>
--	--	---

Table 6: Assembly traductions of the branch functions using gcc 14.2 x86-64.

The assembly presented in Table 6 shows the naive version using a conditional check inside the loop, which can lead to performance issues due to branch mispredictions. Moreover, the condition being applied on every iteration is very instruction-heavy (test and je being called every cycle).

The manually optimized function performs the check a single time and then loops normally. This reduces the number of instruction and reads done on each iteration.

The compiler-optimized version recognized that the loop invariant of successfully extracted it. It also unfortunately unrolled the loop and applied shifting instead of multiplication, a side effect of using a generic optimization flag.

DTMF program optimizations

The whole DTMF program is already compiled with `Ofast`, which suggests that the compiler has already applied a lot of optimizations. It was very difficult to find regions with code that was obviously not optimized in the first place.

Cache result in `_dtmf_normalize_signal` function

It was found that the `_dtmf_normalize_signal` function of the `dtmf_common.c` file that computes the same value twice instead of caching it.

```
1  static void _dtmf_normalize_signal(dtmf_float_t *buffer, dtmf_count_t const
count) {
2      dtmf_float_t max = 0.0;
3      for (dtmf_count_t i = 0; i < count; i++) {
4          if (fabs(buffer[i]) > max)
5              max = fabs(buffer[i]);
6      }
7      ...
8  }
```

It can be optimized by caching the result of `fabs(buffer[i])` in a local variable, which can be reused in the comparison. This reduces the number of calls to `fabs` and improves performance.

```
1  static void _dtmf_normalize_signal(dtmf_float_t *buffer, dtmf_count_t const
count) {
2      dtmf_float_t max = 0.0;
3      // Find maximum absolute value
4      for (dtmf_count_t i = 0; i < count; i++) {
5          dtmf_float_t abs_val = fabs(buffer[i]);
6          if (abs_val > max)
7              max = abs_val;
8      }
9      ...
10 }
```

Without listing the assembly here, it is cleaner and the compiler indeed makes a single call to `fabs`. The assembly can be seen online at <https://godbolt.org/z/7dxreod9q>.

Conclusion

In conclusion, we observe that the compiler is able to apply a lot of optimizations automatically. However, it is not always clear which flags are needed to apply a specific optimization. The optimizations we have seen in this report are not exhaustive and there are many more flags that can be applied but the time needed to go through the documentation and figure out which is which is extensive.